

Assignment 2

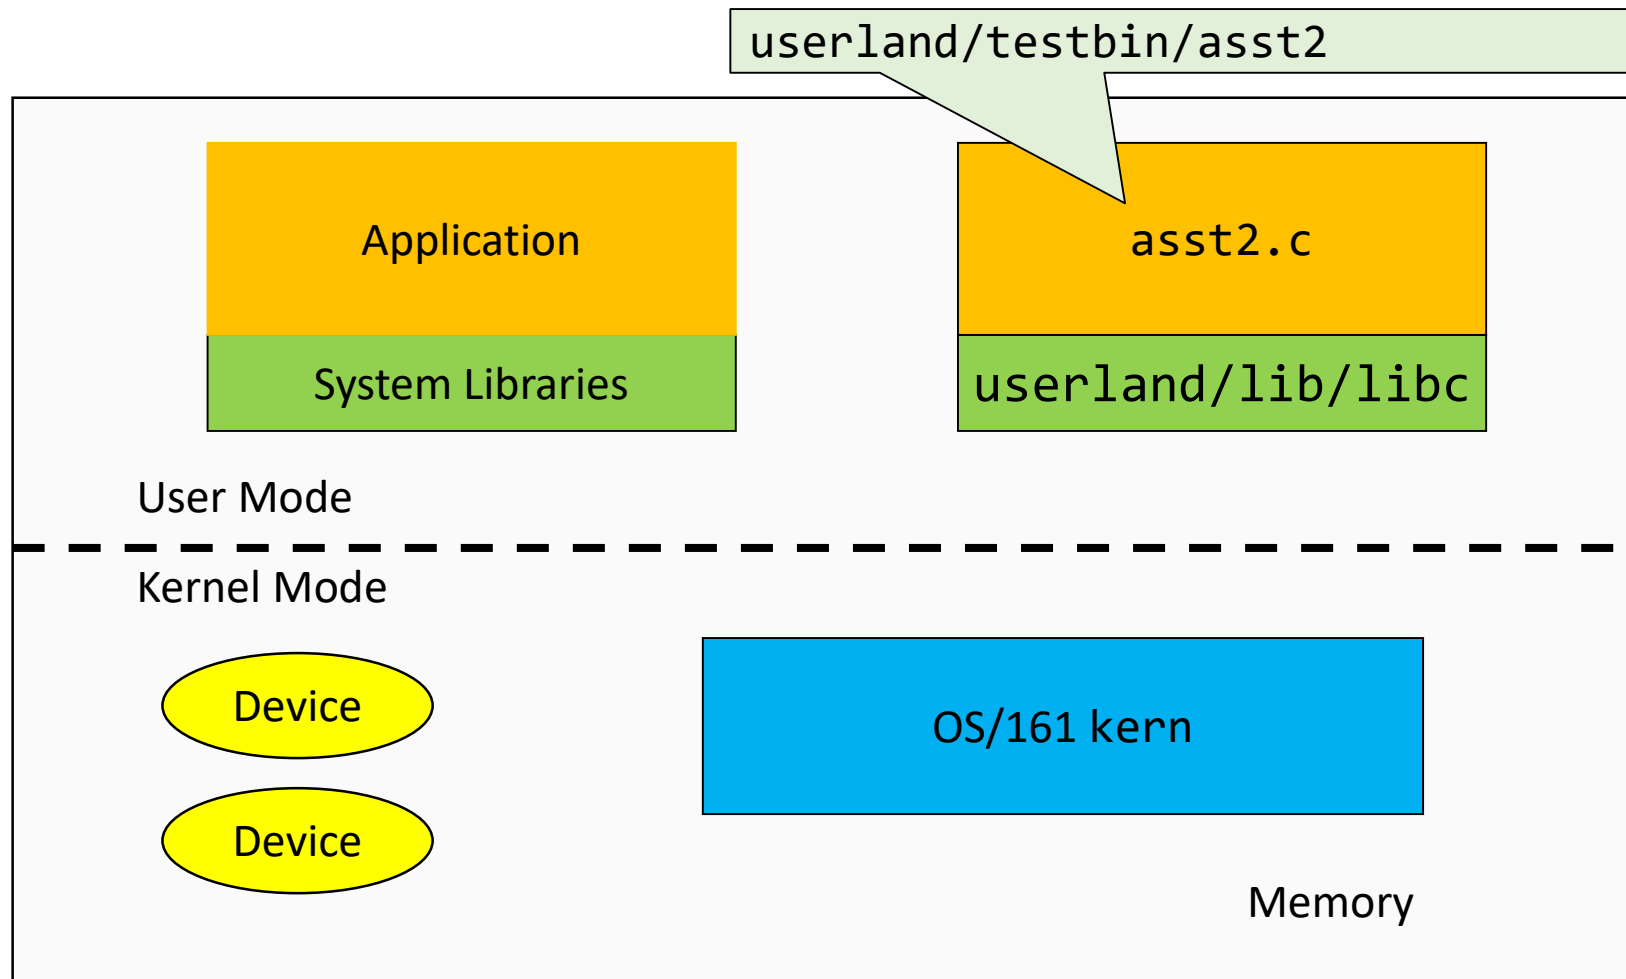
The Basic ASST2 Spec

- Implement `open()`, `read()`, `write()`, `lseek()`, `close()`, and `dup2()`
 - Assume you need to support `fork()`
 - Document the concurrency issues introduced by `fork()`
 - However, you should not synchronise the actual code
 - Can assume we will only test with a single process at a time.
 - Your data structures should not need significant changes to support `fork()`
 - Except for synchronisation
 - User-level exists
 - `asst2`
 - C libraries
 - An existing framework and code for:
 - system call dispatching,
 - VFS
 - Emufs
 - drivers

Overview


- Overall structure
 - User-level
 - Process structure
 - In-kernel
 - The storage stack
 - Overview of VFS and emufs functionality
- Details
 - Understanding the system interface
 - Argument passing
 - System call dispatching
 - Moving data across the user-kernel boundary
 - Connecting the interface to the VFS

Structure of a Computer System



R3000 Address Space Layout

- ksegX not accessible in usermode
- Switching processes switches the application view of memory (translation stored in a page table) for kuseg



Process Layout

- Where is asst2 code/data (from asst2.c)?

0xffffffff

0xC0000000

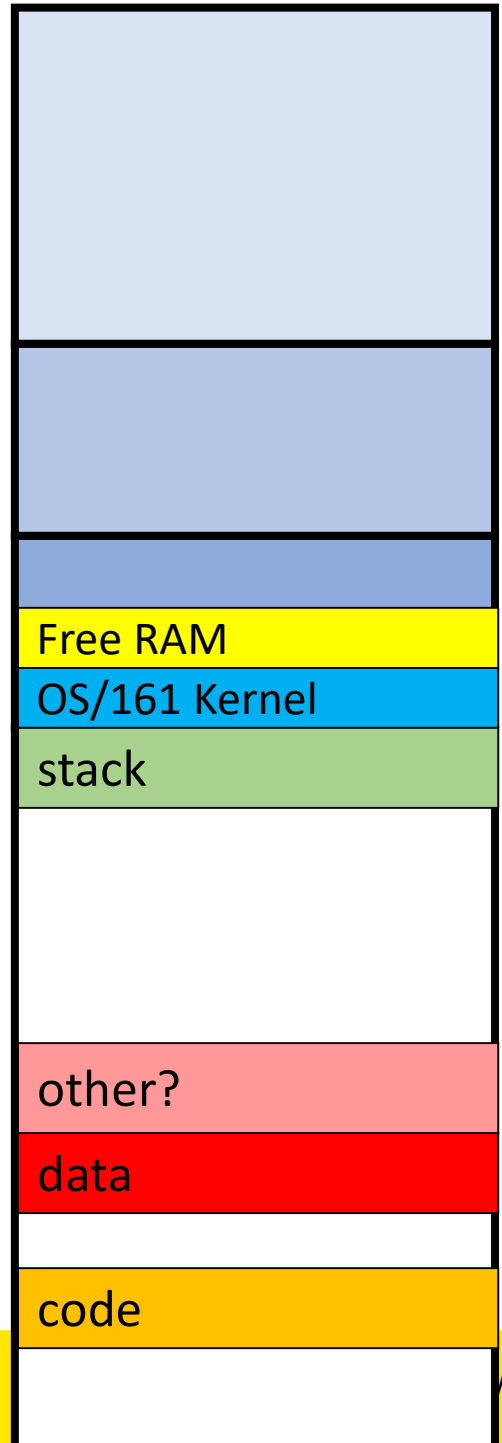
0xA0000000

0x80000000

0x10000000

0x04000000

0x00000000



Calling open()

```
int open(const char *filename,  
        int flags, ...);
```

- Where is the function “open()”?

0xffffffff

0xC0000000


0xA0000000

0x80000000


0x10000000

0x04000000

0x00000000



Structure of a Computer System



OS/161 storage stack

asst2

Application



Syscall dispatching

FD table

OF table

VFS

FS

Device driver

emufs

sys161



~/cs3231/root



open()?

```
int open(const char *filename,  
        int flags, ...);
```

- Where is “open()’s” implementation?
- By convention, it’s called `sys_open()` in the kernel.

This is what you are
implementing in ASST2

0xffffffff

0xC0000000


0xA0000000

0x80000000

0x10000000

0x04000000

0x00000000



Existing storage stack

asst2

Application



Syscall dispatching

FD table

OF table

VFS

FS

Device driver

emufs

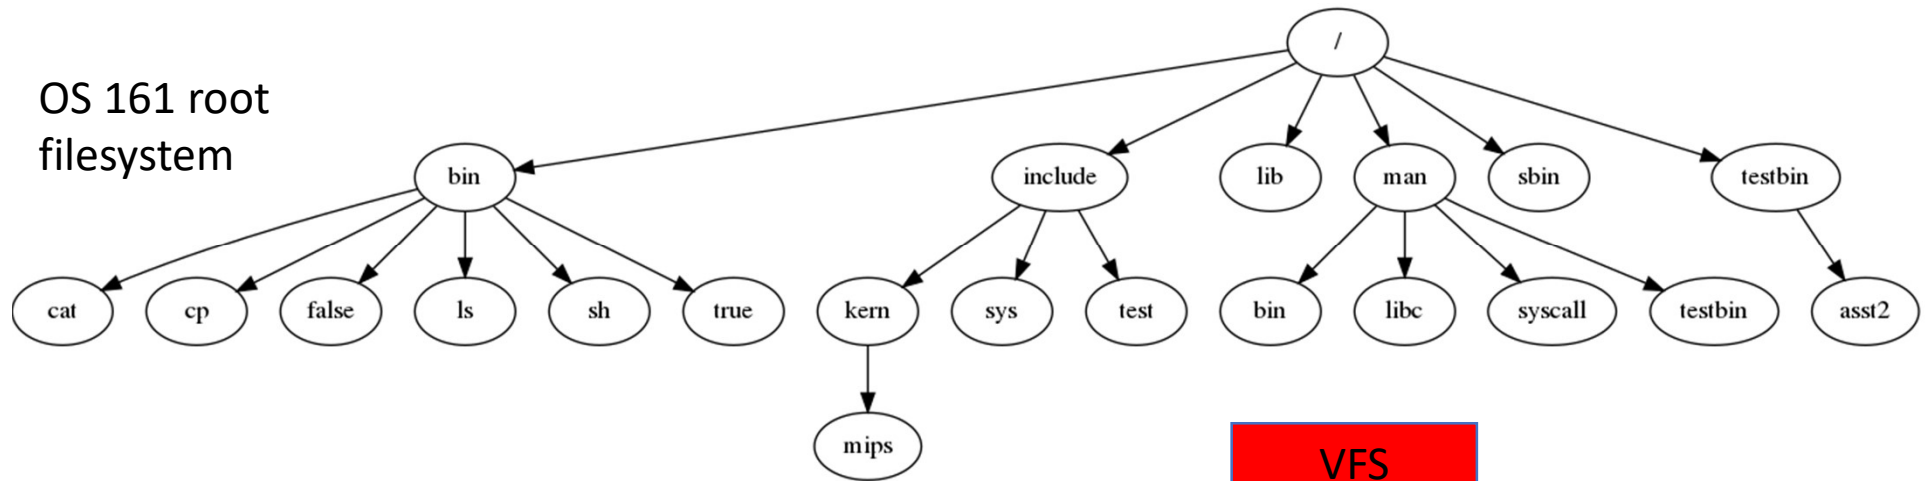
sys161



~/cs3231/root



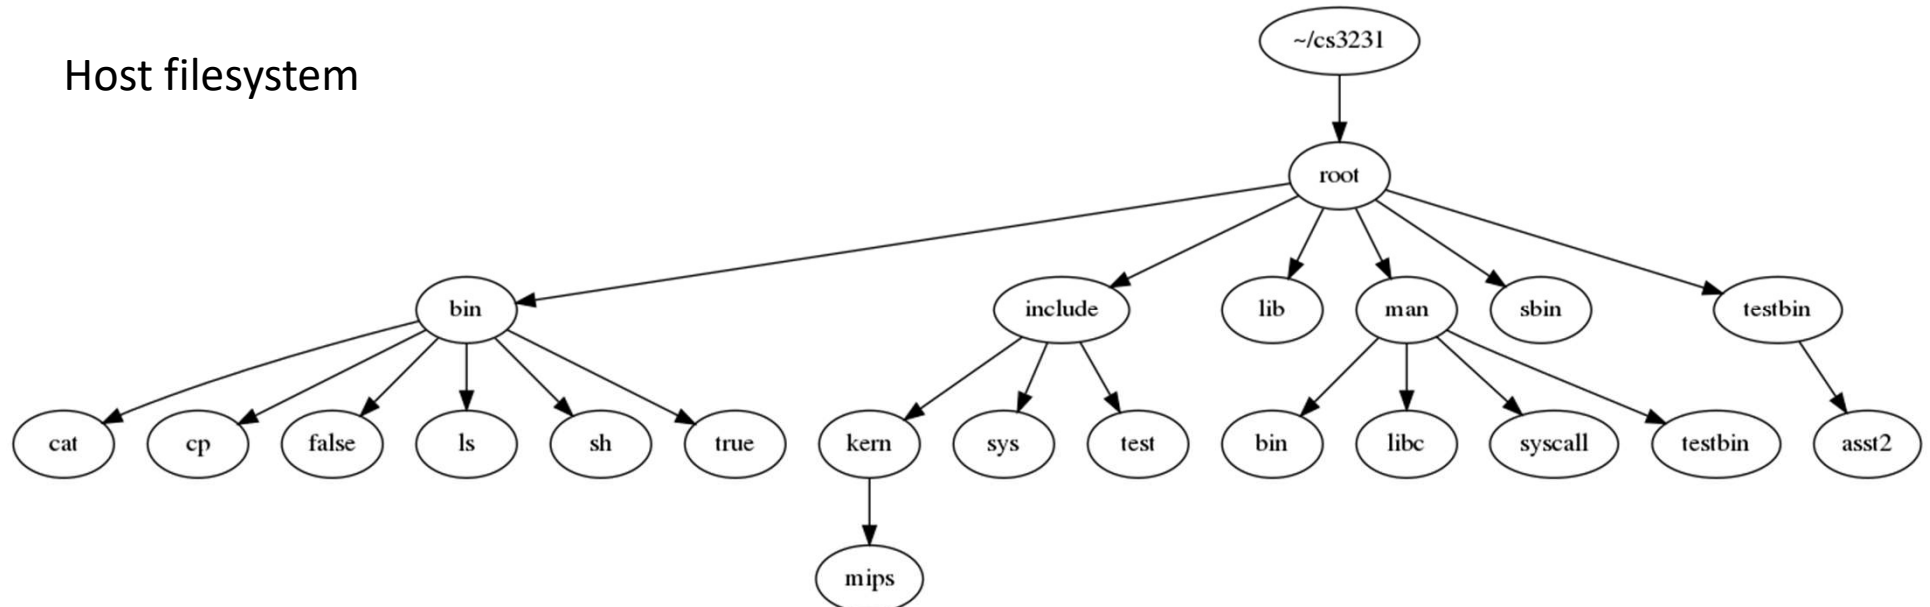
OS 161 root
filesystem



VFS
emufs
driver
sys161

Provided Storage Stack

Host filesystem



Details

System Call Interface

```
int open(const char *filename, int flags);  
int open(const char *filename, int flags, mode_t mode);  
int close(int fd);  
ssize_t read(int fd, void *buf, size_t buflen);  
ssize_t write(int fd, const void *buf, size_t nbytes);  
int dup2(int oldfd, int newfd);  
off_t lseek(int fd, off_t pos, int whence);
```

Solution should work with fork() if implemented

```
pid_t fork(void);
```

open/close

```
int open(const char *filename, int flags);  
int open(const char *filename, int flags, mode_t mode);  
int close(int fd);
```

Read/write

```
ssize_t read(int fd, void *buf, size_t buflen);  
ssize_t write(int fd, const void *buf, size_t nbytes);
```


dup2

```
int dup2(int oldfd, int newfd);
```

lseek

```
off_t lseek(int fd, off_t pos, int whence);
```

fork

```
pid_t fork(void);
```

Argument passing

```
#include <unistd.h>
```

```
int reboot(int code);
```

Description

reboot reboots or shuts down the system. The specific action depends on the code passed:

RB_REBOOT The system is rebooted.

RB_HALT The system is halted.

RB_POWEROFF The system is powered off.

Return Values

On success, reboot does not return. On error, -1 is returned, and errno is set according to the error encountered.

Convention for
kernel entry

| |
|------|
| ra |
| fp |
| sp |
| gp |
| k1 |
| k0 |
| s7 |
| ⋮ |
| s0 |
| t9 |
| ⋮ |
| t0 |
| a3 |
| a2 |
| a1 |
| a0 |
| v1 |
| v0 |
| AT |
| zero |

Preserved



Preserved for C
calling
convention

Preserved



Args in



Success?



Result



SysCall No.

Convention
for kernel
exit

| |
|------|
| ra |
| fp |
| sp |
| gp |
| k1 |
| k0 |
| s7 |
| ⋮ |
| s0 |
| t9 |
| ⋮ |
| t0 |
| a3 |
| a2 |
| a1 |
| a0 |
| v1 |
| v0 |
| AT |
| zero |

```

struct trapframe {
    u_int32_t tf_vaddr; /* vaddr register */
    u_int32_t tf_status; /* status register */
    u_int32_t tf_cause; /* cause register */
    u_int32_t tf_lo;
    u_int32_t tf_hi;
    u_int32_t tf_ra; /* Saved register 31 */
    u_int32_t tf_at; /* Saved register 1 (AT) */
    u_int32_t tf_v0; /* Saved register 2 (v0) */
    u_int32_t tf_v1; /* etc. */
    u_int32_t tf_a0;
    u_int32_t tf_a1;
    u_int32_t tf_a2;
    u_int32_t tf_a3;
    u_int32_t tf_t0;
    ...
    u_int32_t tf_t7;
    u_int32_t tf_s0;
    ...
    u_int32_t tf_s7;
    u_int32_t tf_t8;
    u_int32_t tf_t9;
    u_int32_t tf_k0;
    */
    u_int32_t tf_k1;
    u_int32_t tf_gp;
    u_int32_t tf_sp;
    u_int32_t tf_s8;
    u_int32_t tf_epc; /* coprocessor 0 epc register */
};

```

By creating a pointer to here of type struct trapframe *, we can access the user's saved registers as normal variables within 'C'

Kernel Stack

| |
|--------|
| |
| epc |
| s8 |
| sp |
| gp |
| k1 |
| k0 |
| t9 |
| t8 |
| ⋮ |
| at |
| ra |
| hi |
| lo |
| cause |
| status |
| vaddr |

```
syscall(struct trapframe *tf)
{
    callno = tf->tf_v0;
    retval = 0;

    switch (callno) {
        case SYS_reboot:
            err = sys_reboot(tf->tf_a0);
            break;

        /* Add stuff here */

        default:
            kprintf("Unknown syscall %d\n", callno);
            err = ENOSYS;
            break;
    }
}
```

```
if (err) {
    tf->tf_v0 = err;
    tf->tf_a3 = 1;      /* signal an error */
}
else {
    /* Success. */
    tf->tf_v0 = retval;
    tf->tf_a3 = 0;      /* signal no error */
}

tf->tf_epc += 4;

}
```


System Call Interface

```
int open(const char *filename, int flags);  
int open(const char *filename, int flags, mode_t mode);  
int close(int fd);  
ssize_t read(int fd, void *buf, size_t buflen);  
ssize_t write(int fd, const void *buf, size_t nbytes);  
int dup2(int oldfd, int newfd);  
off_t lseek(int fd, off_t pos, int whence);
```

lseek() Offset

```
uint64_t offset;
```

```
int whence;
```

```
off_t retval64;
```

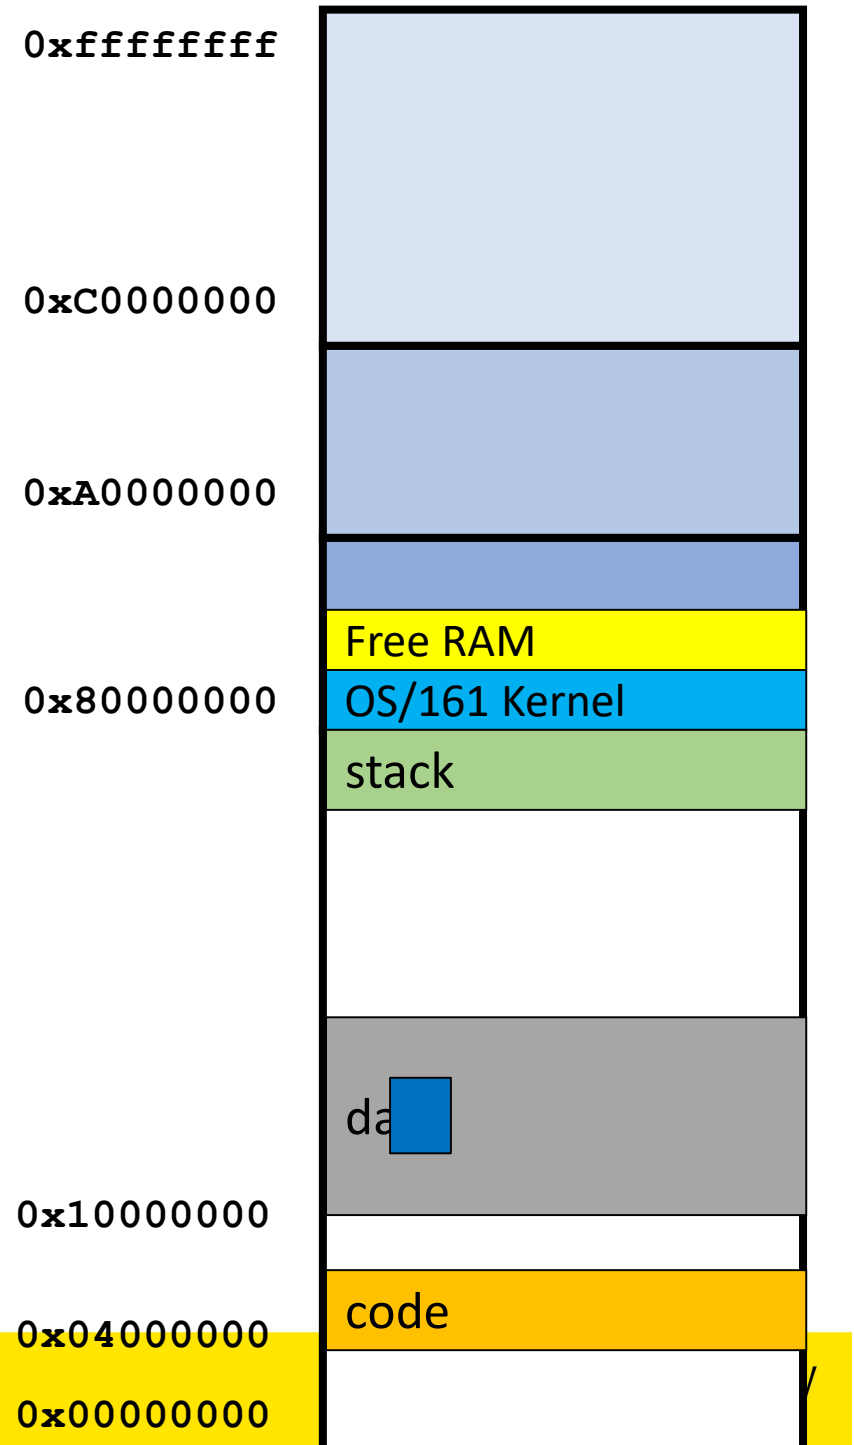
```
join32to64(tf->tf_a2, tf->tf_a3, &offset);
```

```
copyin((userptr_t)tf->tf_sp + 16, &whence, sizeof(int));
```

```
split64to32(retval64, &tf->tf_v0, &tf->tf_v1);
```


Pointers

- What about the first argument to `open()`
 - It's a string?
- What are the problems with accessing a string (i.e. user-specified region of memory)?



Copy in/out(str)

```
int copyin(const_userptr_t usersrc, void *dest,  
           size_t len);  
int copyout(const void *src, userptr_t userdest,  
            size_t len);  
int copyinstr(const_userptr_t usersrc, char  
              *dest, size_t len, size_t *got);  
int copyoutstr(const char *src, userptr_t  
               userdest, size_t len, size_t *got);
```



0xffffffff

0xC0000000

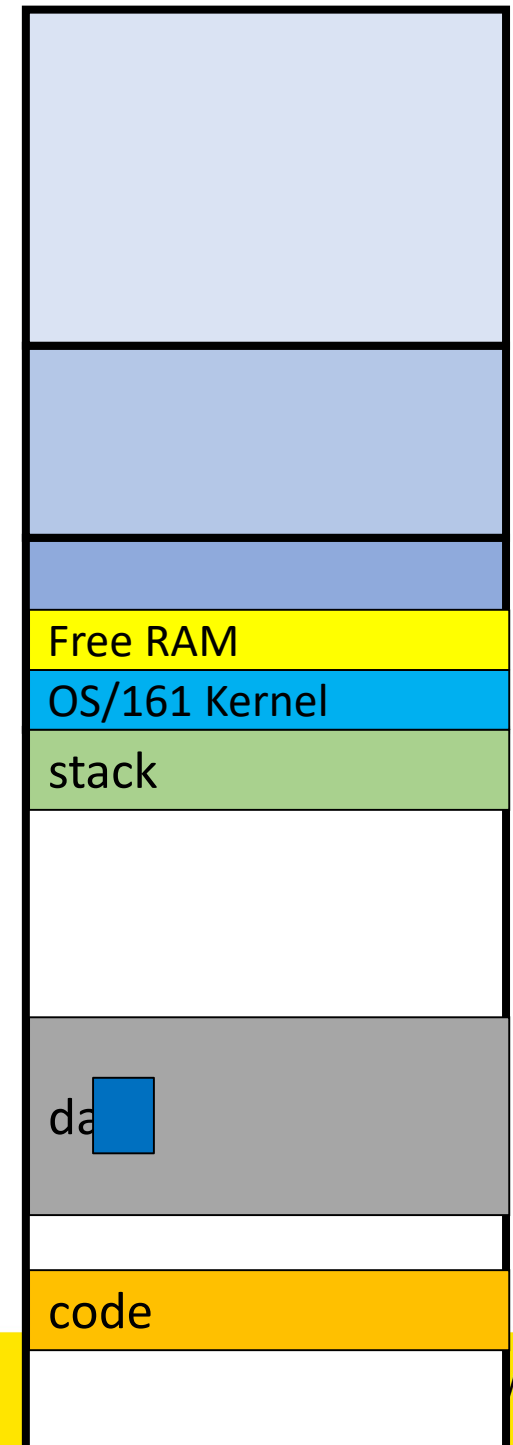
0xA0000000

0x80000000

0x10000000

0x04000000

0x00000000



Buffers – e.g. read()

- Kernel framework for safely handling buffers
 - Does error/range/validity checking for you

```
ssize_t read(int fd, void *buf, size_t buflen);
```

```
struct iovec {  
    union {  
        userptr_t  iov_ubase;    /* user-supplied pointer */  
        void        *iov_kbase;  /* kernel-supplied pointer */  
    };  
    size_t iov_len;  /* Length of data */  
};
```

0xffffffff

0xc0000000


0xa0000000

0x80000000

0x10000000

0x04000000

0x00000000



VFS READ

A macro with sanity checking

`VOP_READ(vn, uio)`

Invokes a function point of following prototype:

```
int (*vop_read)(struct vnode *file, struct uio *uio);
```

What are the arguments?

UIO

```
/* Source/destination. */
```

```
enum uio_seg {
```

```
    UIO_USERISPACE,
```

```
    /* User process code. */
```

```
    UIO_USERSPACE,
```

```
    /* User process data. */
```

```
    UIO_SYSSPACE,
```

```
    /* Kernel. */
```

```
};
```



```
struct uio {
```

```
    struct iovec    *uio_iov;    /* Data blocks */
```

```
    unsigned        uio_iovcnt; /* Number of iovecs */
```

```
    off_t           uio_offset; /* Desired offset into object */
```

```
    size_t          uio_resid; /* Remaining amt of data to xfer */
```

```
    enum uio_seg     uio_segflg; /* What kind of pointer we have */
```

```
    enum uio_rw      uio_rw; /* Whether op is a read or write */
```

```
    struct addrspace *uio_space; /* Address space for user pointer */
```

```
};
```


Sample Helper function

```
uio_uinit(struct iovec *iov, struct uio *u, userptr_t buf,
size_t len, off_t offset, enum uio_rw rw)
{
    iov->iov_ubase = buf;
    iov->iov_len = len;
    u->uio_iov = iov;
    u->uio_iovcnt = 1;
    u->uio_offset = offset;
    u->uio_resid = len;
    u->uio_segflg = UIO_USERSPACE;
    u->uio_rw = rw;
    u->uio_space = proc_getas();
}
```

System call implementation

1. `sys_open()`
2. `sys_close()`
3. `sys_read()`
4. `sys_write()`
5. `sys_lseek()`
6. `sys_dup2()`

1. `vfs_open()`
 - `copyinstr()`
2. `vfs_close()`
3. `VOP_READ()`
4. `VOP_WRITE()`
5. `VOP_ISSEEKABLE()`
6. `VOP_STAT()`